

MP2-Implementation

r09922136 廖婕吟

1. VMA management in process's user address space

1) VMA structure is defined in kernel/proc.h

```

1 // Virtual Memory area
2 struct vma {
3     struct spinlock lock;
4
5     // p->lock must be held when using these:
6     uint64 start;        // VMA start, inclusive
7     uint64 end;          // VMA end , exclusive
8     struct vma *next;    // list of VMA's
9     uint64 length;       // 0 means vma not used
10    uint64 off;           // offset within file
11    int page_prot;        // access permissions
12    int flags;            // shared or private
13    struct file *file;    // mapped file, if any
14 };

```

A fixed-size array `vma_list[]` in length 16 of VMAs is declared in `kernel/proc.c`, and we allocate VMA from this array as needed.

2) VMAs is allocate from the middle of the process's user address space since defining `VMA_START` at the position of `MAXVA / 2`, and the direction we allocate the VMA is from the lower address to the higher address.

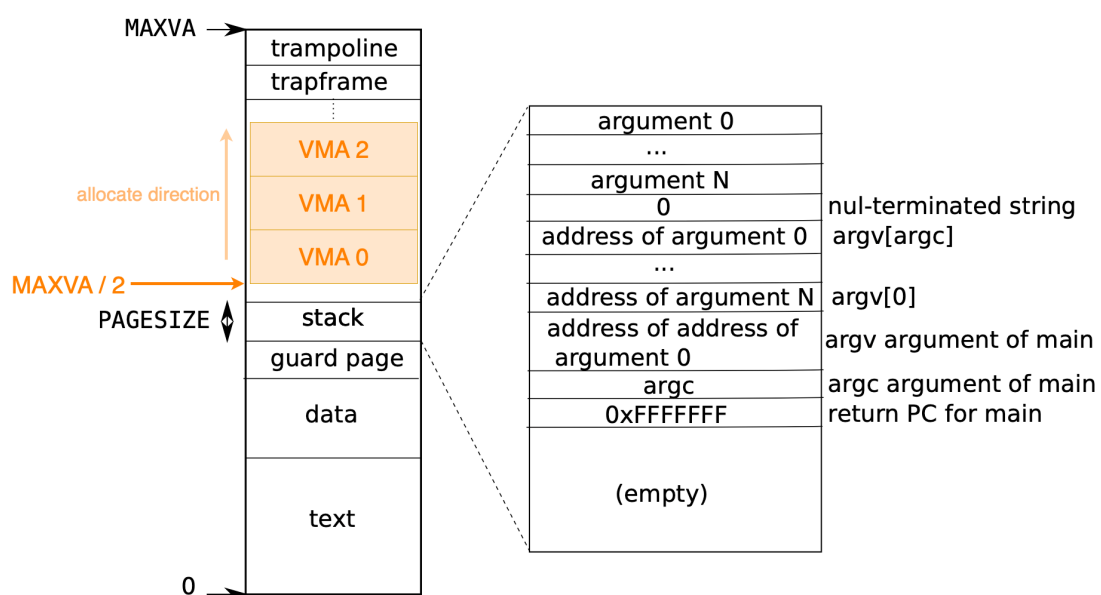


Figure 1: Allocate VMAs from the middle of a process's user address space

- 3) The address that we start to allocate a VMA (`uint64 start`) which is declare in the `struct vma` is always round to a multiple of the page size. that is we allocate VMA with page-aligned.




Figure 2: Allocate VMA with page aligned

- 4) The memory area that we suppose to `sys_munmap()` is only can be unmap from the start or from the end of a VMA, which means we don't support splitting a VMA into two VMAs when unmapping.
- 5) Use link list data structure to maintain the **process's table of mapped regions**. When `sys_mmap()` is called, the new VMA will be added into the list of VMAs by using the vma pointer (`struct vma *next`) which declair in the `struct vma`.

2. Assumptions

- 1) `addr` will always be zero.
- 2) `offset` will always be zero, that is, we always map the file from the starting point of the `mmap`-ed file.
- 3) When `length` is 0, the VMA is not used.
- 4) `flags` will be one of the following bits
 - `MAP_SHARED`: modifications to the mapped memory should be written back to the file.
 - `MAP_PRIVATE`: modifications should not be written back.
- 5) The VMA is page-aligned when we `sys_mmap()`.
- 6) the `sys_munmap()` is not supported to unmap the area that is at the middle of a VMA, that is only can `munmap` from the start or from the end of a VMA and won't split a VMA into two VMAs.